

## John Barker, Programming with R Program Assignment 2

```
##### (1) R has probability functions available for use (Kabacoff, Section 5.2.3).  
# Using one distribution to approximate another is not uncommon.  
# (1)(a) (6 points) The Poisson distribution may be used to approximate the binomial distribution  
# if n > 20 and np < 7. Estimate the following binomial probabilities using  
# *dpois()* and *ppois()* with probability p = 0.05, and n = 100. Then, estimate  
# the same probabilities using *dbinom()* and *pbinom()*.  
# Show the numerical results of your calculations.
```

```
# (i) The probability of exactly 0 successes.
```

```
set.seed(1234)
```

```
dpois(0,5)
```

```
ppois(0,5)
```

```
dbinom(0,100,.05)
```

```
pbinom(0, 100,.05)
```

```
output
```

```
> set.seed(1234)  
> dpois(0,5)  
[1] 0.006737947  
> ppois(0,5)  
[1] 0.006737947  
>  
> dbinom(0,100,.05)  
[1] 0.005920529  
> pbinom(0, 100,.05)  
[1] 0.005920529
```

#(ii) The probability of fewer than 6 successes.

```
set.seed(1234)  
dpois(5,5)  
ppois(5,5)-ppois(4,5)  
dbinom(5,100,.05)  
pbinom(5,100,.05)-pbinom(4,100,.05)
```

output

```
-----  
> set.seed(1234)  
> dpois(5,5)  
[1] 0.1754674  
> ppois(5,5)-ppois(4,5)  
[1] 0.1754674  
> dbinom(5,100,.05)  
[1] 0.1800178  
> pbinom(5,100,.05)-pbinom(4,100,.05)  
[1] 0.1800178
```

```
# (1)(b) (3 points) Generate side-by-side barplots using *par(mfrow = c(1,2))* or  
# *grid.arrange()*. The left barplot will show Poisson probabilities for outcomes  
# ranging from 0 to 10. The right barplot will show binomial probabilities for  
# outcomes ranging from 0 to 10. Use p = 0.05 and n = 100. Title each plot,  
# present in color and assign names to the bar; i.e. x-axis value labels.
```

```
n=100  
p=.05  
lambda<-n*p  
trials <- c(0:10)  
successes<-factor(trials)  
pprobabilities <- dpois(trials,lambda = lambda)  
bprobabilities <- dbinom(trials,size=n,prob=p)
```

```

par(mfrow=c(1,2))

barplot(pprobabilities, names.arg = successes, xlab = "successes", ylab = "Poisson probabilities",
       main = "Poisson Probabilities trials = 10, n= 100 p= .05", col="red")

barplot(bprobabilities, names.arg = successes, xlab = "successes", ylab = "binomial probabilities",
       main = "Binomial Probabilities trials = 10, n= 100 p= .05", col = "darkblue")

par(mfrow=c(1,1))

```

output

---

```

n=100

> p=.05

> lambda<-n*p

> trials <- c(0:10)

> successes<-factor(trials)

>

>

> pprobabilities <- dpois(trials,lambda = lambda)

> bprobabilities <- dbinom(trials,size=n,prob=p)

>

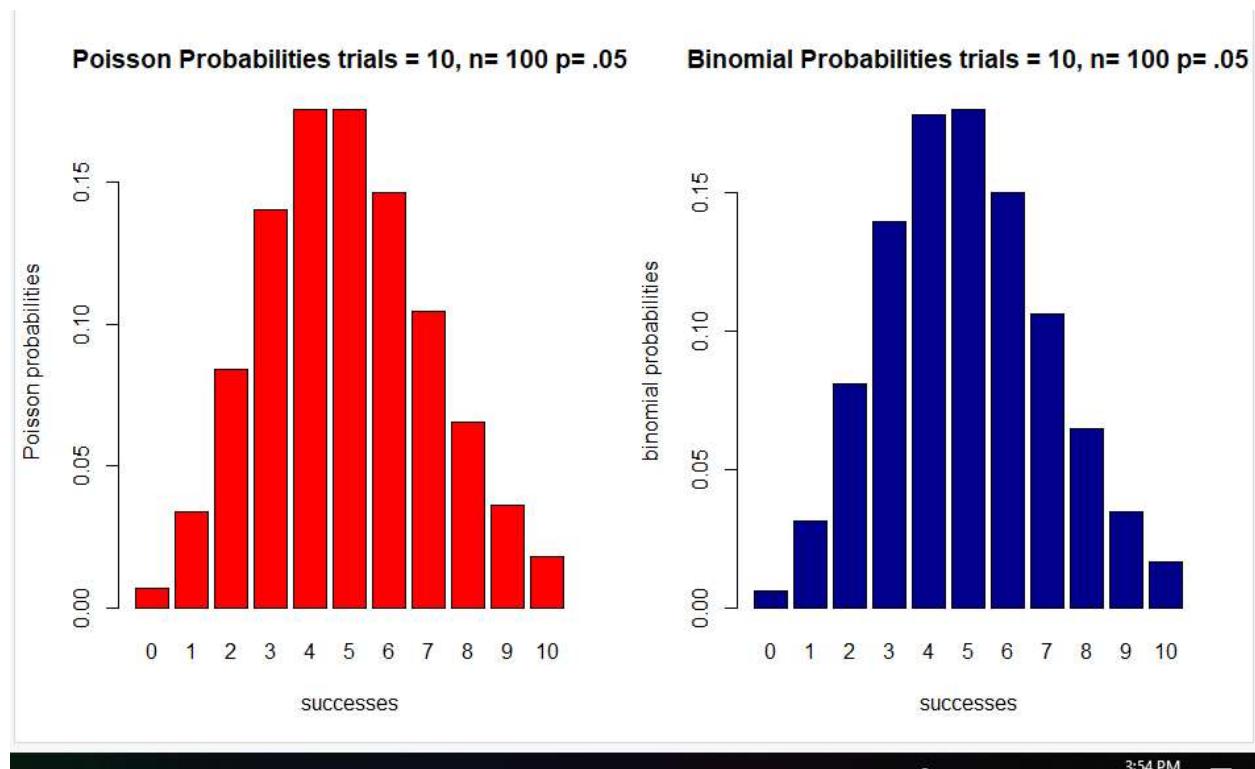
> par(mfrow=c(1,2))

> barplot(pprobabilities, names.arg = successes, xlab = "successes", ylab = "Poisson probabilities",
+       main = "Poisson Probabilities trials = 10, n= 100 p= .05", col="red")

> barplot(bprobabilities, names.arg = successes, xlab = "successes", ylab = "binomial probabilities",
+       main = "Binomial Probabilities trials = 10, n= 100 p= .05", col = "darkblue")

> par(mfrow=c(1,1))

```



# (1)(c) For this problem, refer to Sections 5.2 of Business Statistics. A

# discrete random variable has outcomes: 0, 1, 2, 3, 4, 5, 6. The

# corresponding probabilities in sequence with the outcomes are:

# 0.215, 0.230, 0.240, 0.182, 0.130, 0.003, 0.001. In other words,

# the probability of obtaining "0" is 0.215.

# (i) (3 points) Calculate the expected value and variance for this distribution

# using the general formula for mean and variance of a discrete distribution. To do this,

# you will need to use integer values from 0 to 6 as outcomes along with the

# corresponding probabilities. Round your answer to 2 decimal places.

```
discretes<-c(0,1,2,3,4,5,6)
```

```
discprobs<-c(0.215, 0.230, 0.240, 0.182, 0.130, 0.003, 0.001)
```

```
discmean<round(sum(discretes*discprobs),digits=2)
```

```
discmean
```

```
discvar<-round(sum(((discretes-discmean)**2)*discprobs),digits=2)
discvar

output
#####
> discretes<-c(0,1,2,3,4,5,6)
> discprobs<-c(0.215, 0.230, 0.240, 0.182, 0.130, 0.003, 0.001)
> discmean<-round(sum(discretes*discprobs),digits=2)
> discmean
[1] 1.8
> discvar<-round(sum(((discretes-discmean)**2)*discprobs),digits=2)
> discvar
[1] 1.79
```

```
# (ii) (3 points) Use the *cumsum()* function and plot the cumulative
# probabilties versus the corresponding outcomes. Detemine the value
# of the median for this distribution and show on this plot.
```

```
disccum<-c(cumsum(discprobs))
```

```
disccum
```

```
index<-0
```

```
median<-0
```

```
median
```

```
plot(discretes, disccum,
      main="Cumulative Probability",
      xlab="Discrete Values",
      ylab="Cumulative Probability")
```

```
lines(discretes, disccum)
```

```
abline(v=2)
```

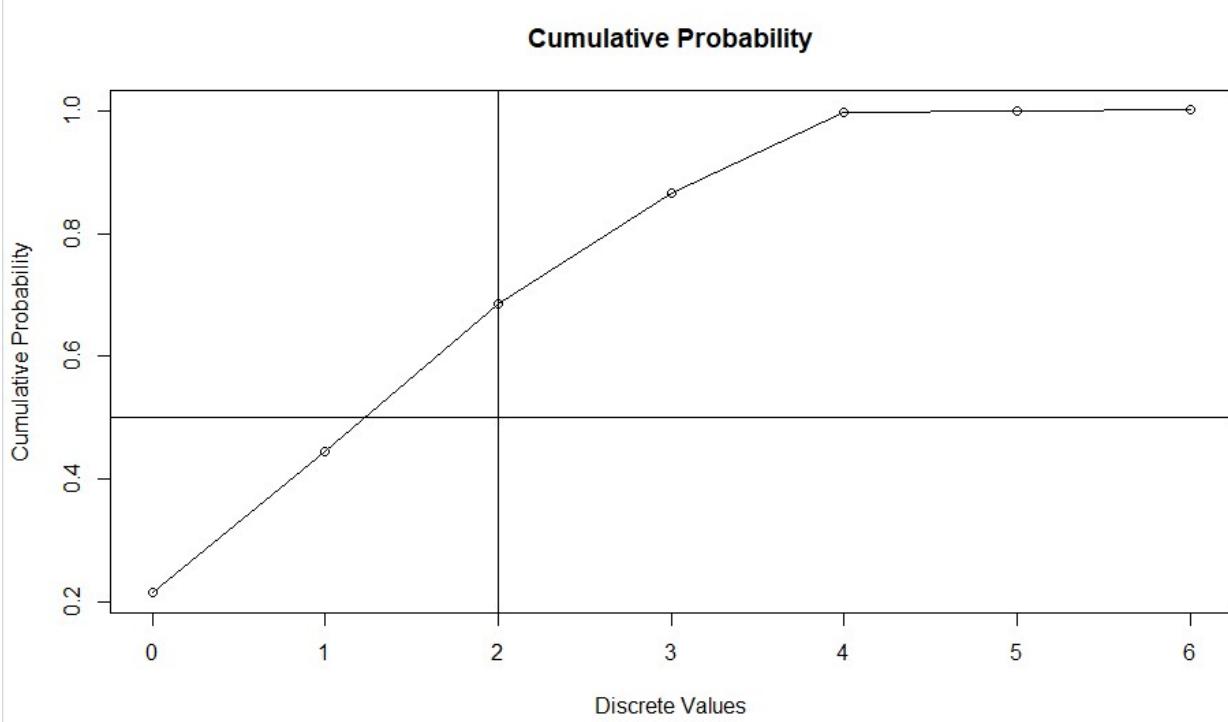
```
abline(h=.5)
```

```
disccum
```

```

output
-----
> disccum<-c(cumsum(discprobs))
> disccum
[1] 0.215 0.445 0.685 0.867 0.997 1.000 1.001
> index<-0
> median<-0
> median
[1] 0
> plot(discretes, disccum,
+       main="Cumulative Probability",
+       xlab="Discrete Values",
+       ylab="Cumulative Probability")
> lines(discretes, disccum)
> abline(v=2)
> abline(h=.5)
> disccum
[1] 0.215 0.445 0.685 0.867 0.997 1.000 1.001

```



##### Section 2: (15 points) #####

##### (2) Conditional probabilities appear in many contexts and, in particular,  
# are used by Bayes' Theorem. Correlations are another means for evaluating dependency  
# between variables. The dataset "faithful"" is part of the "datasets" package  
# and may be loaded with the statement \*data(faithful)\*. It contains 272 observations  
# of 2 variables; waiting time between eruptions (in minutes) and the duration of the

```

# eruption (in minutes) for the Old Faithful geyser in Yellowstone National Park.

#(2)(a) (3 points) Load the "faithful" dataset and present summary statistics and a histogram
# of waiting times. Additionally, compute the empirical conditional probability of an
# eruption less than 3.0 minutes, if the waiting time exceeds 70 minutes.

data(faithful)
x<-faithful
summary(x)
a<-hist(x$waiting)
pa<-sum(x$waiting>70)/nrow(x)
pa
(Nab <- nrow(x[which(x$waiting>70 & x$eruptions<3),]))
(Nb <- sum(hist(x$eruptions,plot=F)$counts[hist(x$eruptions,plot=F)$mids<3]))
(Pab <- Nab / Nb)

Output
#####
> data(faithful)
> x<-faithful
>
> summary(x)
   eruptions      waiting
  Min.   :1.600  Min.   :43.0
  1st Qu.:2.163  1st Qu.:58.0
  Median :4.000  Median :76.0
  Mean   :3.488  Mean   :70.9
  3rd Qu.:4.454  3rd Qu.:82.0
  Max.   :5.100  Max.   :96.0
>
> a<-hist(x$waiting)
> pa<-sum(x$waiting>70)/nrow(x)
> pa
[1] 0.6066176
> (Nab <- nrow(x[which(x$waiting>70 & x$eruptions<3),]))
[1] 1
> (Nb <- sum(hist(x$eruptions,plot=F)$counts[hist(x$eruptions,plot=F)$mids<3]
))
[1] 97
>
> (Pab <- Nab / Nb)
[1] 0.01030928

```

```
#(i) (3 points) Identify any observations in "faithful" for which the waiting time  
# exceeds 70 minutes and the eruptions are less than 3.0 minutes. List and s  
how any  
# such observations in a distinct color on a scatterplot of all eruption (ver  
tical axis)  
# and waiting times (horizontal axis). Include a horizontal line at eruption  
= 3.0,  
# and a vertical line at waiting time = 70. Add a title and appropriate text  
.
```

```
pabt<-subset(x, subset = waiting>=70 & eruptions <= 3)
```

pabt

```
summary(pabt)
```

str(pabt)

```
color = rep(NA, length=length(x))
```

```
color[which(x$eruptions<3 & x$waiting > 70)] = "red"
```

```
color[which(!(x$eruptions<3 & x$waiting > 70))] = "blue"
```

color

```
plot(x$waiting,x$eruptions, col = color, pch=16, xlab = "Wait time for eruption",
```

```
ylab = "Eruption Time", main = "Scatterplot of Wait Times versus Eruption Times")
```

## output

```
> pabt<-subset(x, subset = waiting>=70 & eruptions <= 3)
```

> *pabt*

211 eruptions waiting  
 2.383 71

```
> summary(pabt)
   eruptions      waiting
Min.    :2.383     Min.    :71
1st Qu.:2.383     1st Qu.:71
Median :2.383     Median :71
Mean   :2.383     Mean   :71
3rd Qu.:2.383     3rd Qu.:71
Max.   :2.383     Max.   :71
```

```
> str(pabt)
```

```
'data.frame': 1 obs.  
$ eruptions: num 2.38  
$ waiting : num 71
```

1

```
> color = rep(NA, length=length(x))
```

```
> color[which(x$eruptions<3 & x$waiting > 70)] = "red"
```

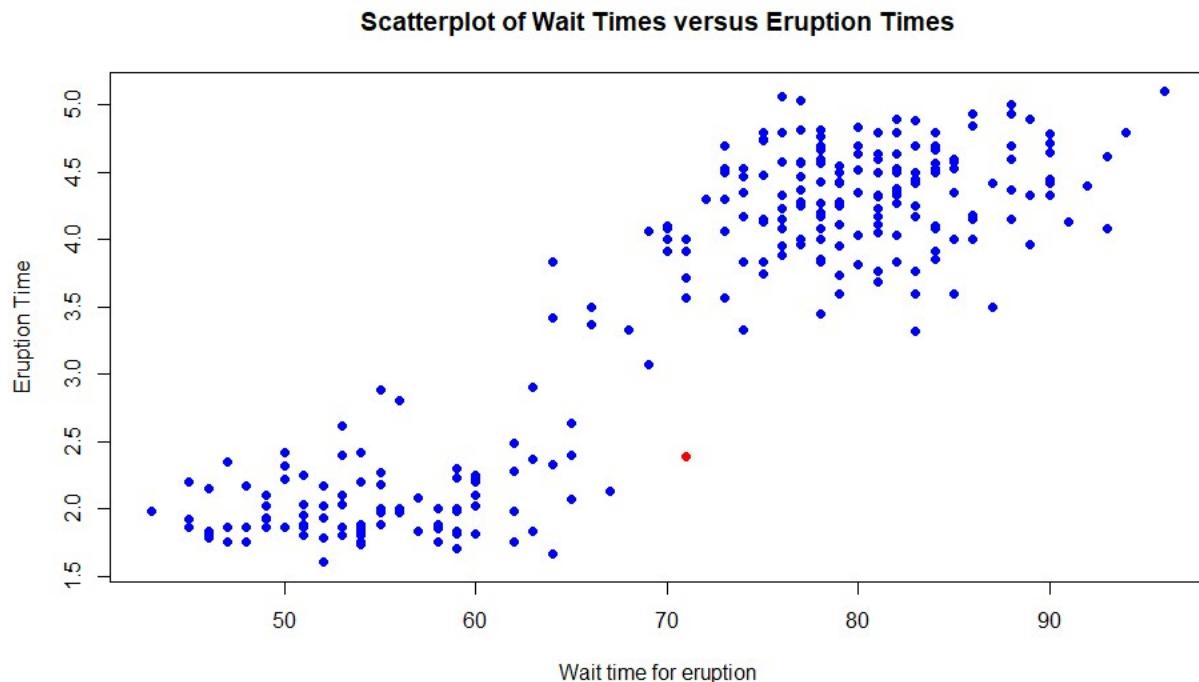
```
> color[which(!(x$eruptions < 3 & x$waiting > 70))] = "blue"
```

> color

```
[1] "blue"  
[16] "blue"  
[31] "blue"  
[46] "blue" "blue" "blue" "blue" "blue" "blue" "blue" "blue" "blue" "blue"
```

```
[61] "blue" "blue"
[76] "blue" "blue"
[91] "blue" "blue"
[106] "blue" "blue"
[121] "blue" "blue"
[136] "blue" "blue"
[151] "blue" "blue"
[166] "blue" "blue"
[181] "blue" "blue"
[196] "blue" "blue"
[211] "red" "blue" "blue"
[226] "blue" "blue"
[241] "blue" "blue"
[256] "blue" "blue"
[271] "blue" "blue"
```

```
> plot(x$waiting,x$eruptions, col = color, pch=16, xlab = "Wait time for eruption",
+       ylab = "Eruption Time", main = "Scatterplot of Wait Times versus Eruption Times")
```



#(ii) (1.5 point) What does the plot suggest about the relationship between eruption time and waiting time?

\*\*\*\*Answer: (Enter your answer here.) \*\*\*

#The plot suggests the longer the wait time for Faithful to erupt the longer the eruption lasts.

#In particular there is only one eruption less than three minutes where the wait time for #eruption exceeds 70.

#-----

```

#(2)(b) (4.5 points) Past research indicates that the waiting times between consecutive
# eruptions are not independent. This problem will check to see if there is evidence
# of this. Form consecutive pairs of waiting times. In other words, pair the first
# and second waiting times, pair the third and fourth waiting times, and so forth.
# There are 136 resulting consecutive pairs of waiting times. Form a data frame with
# the first column containing the first waiting time in a pair and the second column
# with the second waiting time in a pair. Plot the pairs with the second member of
# a pair on the vertical axis and the first member on the horizontal axis.

```

```

# One way to do this is to pass the vector of waiting times - faithful$waiting -
# to *matrix()*, specifying 2 columns for our matrix, with values organized by row;
# i.e. byrow = TRUE.

```

```

vwaiting<-matrix(x$waiting, ncol = 2, byrow = TRUE)
vwaiting
plot(vwaiting, col="orange", xlab = "Waiting time of First Paired Eruption",
      ylab = "Waiting time of Second Paired Eruption", pch=16,
      main = "Scatter Plot of Eruption Wait Times\nfor Consecutive Pairs of Eruptions ")

```

output

---

```

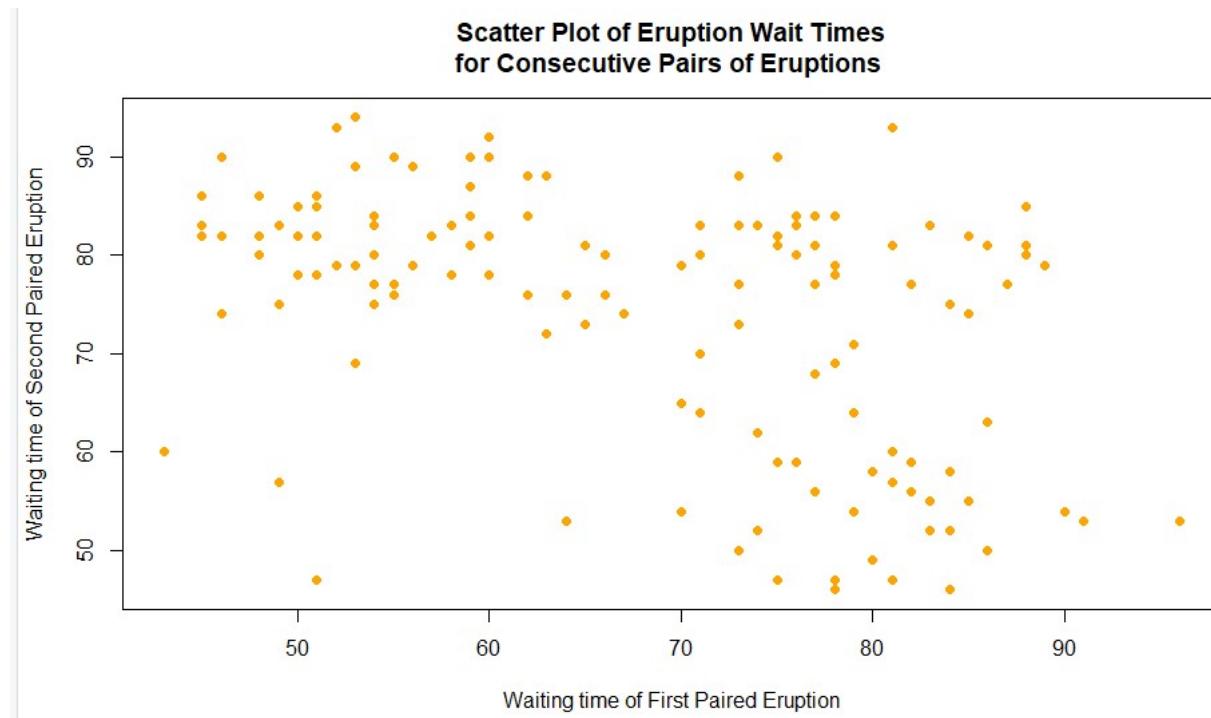
> vwaiting<-matrix(x$waiting, ncol = 2, byrow = TRUE)
> vwaiting
[,1] [,2]
[1,] 79 54
[2,] 74 62
[3,] 85 55
[4,] 88 85
[5,] 51 85
[6,] 54 84
[7,] 78 47
[8,] 83 52
[9,] 62 84
[10,] 52 79
[11,] 51 47
[12,] 78 69
[13,] 74 83
[14,] 55 76
[15,] 78 79
[16,] 73 77

```

[17, ]	66	80
[18, ]	74	52
[19, ]	48	80
[20, ]	59	90
[21, ]	80	58
[22, ]	84	58
[23, ]	73	83
[24, ]	64	53
[25, ]	82	59
[26, ]	75	90
[27, ]	54	80
[28, ]	54	83
[29, ]	71	64
[30, ]	77	81
[31, ]	59	84
[32, ]	48	82
[33, ]	60	92
[34, ]	78	78
[35, ]	65	73
[36, ]	82	56
[37, ]	79	71
[38, ]	62	76
[39, ]	60	78
[40, ]	76	83
[41, ]	75	82
[42, ]	70	65
[43, ]	73	88
[44, ]	76	80
[45, ]	48	86
[46, ]	60	90
[47, ]	50	78
[48, ]	63	72
[49, ]	84	75
[50, ]	51	82
[51, ]	62	88
[52, ]	49	83
[53, ]	81	47
[54, ]	84	52
[55, ]	86	81
[56, ]	75	59
[57, ]	89	79
[58, ]	59	81
[59, ]	50	85
[60, ]	59	87
[61, ]	53	69
[62, ]	77	56
[63, ]	88	81
[64, ]	45	82
[65, ]	55	90
[66, ]	45	83
[67, ]	56	89
[68, ]	46	82
[69, ]	51	86
[70, ]	53	79
[71, ]	81	60
[72, ]	82	77
[73, ]	76	59

[74, ]	80	49
[75, ]	96	53
[76, ]	77	77
[77, ]	65	81
[78, ]	71	70
[79, ]	81	93
[80, ]	53	89
[81, ]	45	86
[82, ]	58	78
[83, ]	66	76
[84, ]	63	88
[85, ]	52	93
[86, ]	49	57
[87, ]	77	68
[88, ]	81	81
[89, ]	73	50
[90, ]	85	74
[91, ]	55	77
[92, ]	83	83
[93, ]	51	78
[94, ]	84	46
[95, ]	83	55
[96, ]	81	57
[97, ]	76	84
[98, ]	77	81
[99, ]	87	77
[100, ]	51	78
[101, ]	60	82
[102, ]	91	53
[103, ]	78	46
[104, ]	77	84
[105, ]	49	83
[106, ]	71	80
[107, ]	49	75
[108, ]	64	76
[109, ]	53	94
[110, ]	55	76
[111, ]	50	82
[112, ]	54	75
[113, ]	78	79
[114, ]	78	78
[115, ]	70	79
[116, ]	70	54
[117, ]	86	50
[118, ]	90	54
[119, ]	54	77
[120, ]	79	64
[121, ]	75	47
[122, ]	86	63
[123, ]	85	82
[124, ]	57	82
[125, ]	67	74
[126, ]	54	83
[127, ]	73	73
[128, ]	88	80
[129, ]	71	83
[130, ]	56	79

```
[131,] 78 84
[132,] 58 83
[133,] 43 60
[134,] 75 81
[135,] 46 90
[136,] 46 74
> plot(vwaiting, col="orange", xlab = "Waiting time of First Paired Eruption"
,
+       ylab = "Waiting time of Second Paired Eruption", pch=16,
+       main = "Scatter Plot of Eruption Wait Times\nfor Consecutive Pairs of
Eruptions ")
```



#(2)(c) (2) Test the hypothesis of independence with a two-sided test at the 5% level

# using the Kendall correlation coefficient.

```
#``{r test2c}
cor(vwaiting, method="kendall", use = "pairwise")
output
#####
> cor(vwaiting, method="kendall", use = "pairwise")
 [,1]      [,2]
[1,] 1.0000000 -0.2935579
[2,] -0.2935579 1.0000000
```

##### Section 3: (15 points) #####

##### (3) Performing hypothesis tests using random samples is fundamental to  
# statistical inference. The first part of this problem involves comparing two  
# different diets. Using "ChickWeight" data available in the base R, "datasets"  
# package, execute the following code to prepare a data frame for analysis.

#``{r test3}

```
# load "ChickWeight" dataset
data(ChickWeight)

# Create T | F vector indicating observations with Time == 21 and Diet == "1" OR "3"
index <- ChickWeight$Time == 21 & (ChickWeight$Diet == "1" | ChickWeight$Diet == "3")

# Create data frame, "result," with the weight and Diet of those observations with "TRUE" "index"" values
result <- subset(ChickWeight[index, ], select = c(weight, Diet))

# Encode "Diet" as a factor
result$Diet <- factor(result$Diet)
str(result)
```

##### The data frame, "result", has chick weights for two diets, identified as

# diet "1" and "3". Use the data frame, "result," to complete the following item.

#(3)(a) (3 points) Display two side-by-side vertical boxplots using par(mfrow = c(1,2)).

# One boxplot would display diet "1" and the other diet "3".

```
par(mfrow=c(1,2))
Diet1<-result[which(result$Diet==1),]
Diet3<-result[which(result$Diet==3),]
boxplot(Diet1$weight, main = "Chick Weights Diet 1", col = "red")
```

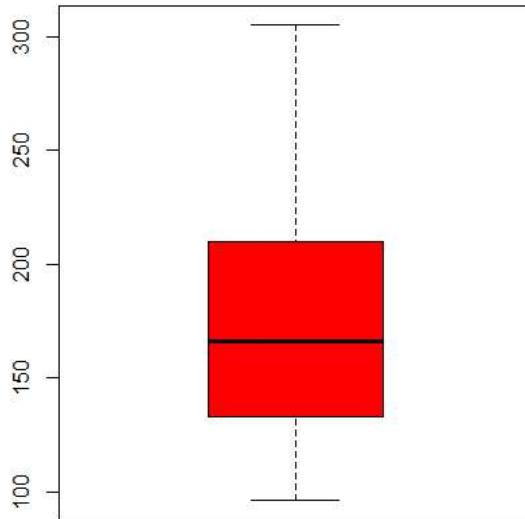
```

boxplot(Diet3$weight, main = "Chick Weights Diet 2", col = "yellow")
par(mfrow=c(1,1))

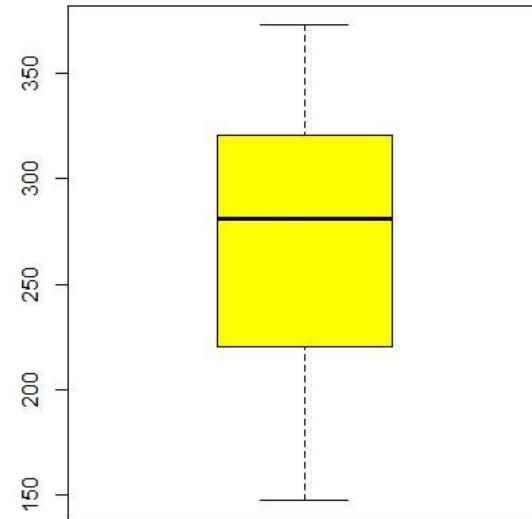
output
#####
> data(ChickWeight)
>
> # Create T | F vector indicating observations with Time == 21 and Diet == "1" OR "3"
> index <- ChickWeight$Time == 21 & (ChickWeight$Diet == "1" | ChickWeight$Diet == "3")
>
> # Create data frame, "result," with the weight and Diet of those observations with "TRUE" "index"" values
> result <- subset(ChickWeight[index, ], select = c(weight, Diet))
>
> # Encode "Diet" as a factor
> result$Diet <- factor(result$Diet)
> str(result)
Classes 'nfnGroupedData', 'nfGroupedData', 'groupedData' and 'data.frame': 2
6 obs. of 2 variables:
 $ weight: num 205 215 202 157 223 157 305 98 124 175 ...
 $ Diet   : Factor w/ 2 levels "1","3": 1 1 1 1 1 1 1 1 1 1 ...
> par(mfrow=c(1,2))
> Diet1<-result[which(result$Diet==1),]
> Diet3<-result[which(result$Diet==3),]
> boxplot(Diet1$weight, main = "Chick weights Diet 1", col = "red")
> boxplot(Diet3$weight, main = "Chick weights Diet 2", col = "yellow")
> par(mfrow=c(1,1))

```

**Chick Weights Diet 1**



**Chick Weights Diet 2**



```
#(3)(b) (3 points) Use the "weight" data for the two diets to test the  
# null hypothesis of equal population mean weights for the two diets. Test at  
# the 95% confidence level with a two-sided t-test. This can be done using  
# *t.test()* in R. Assume equal variances. Display the results of t.test().
```

```
t.test(Diet1$weight, Diet3$weight)
```

output

```
> t.test(Diet1$weight, Diet3$weight)  
  
Welch Two Sample t-test  
  
data: Diet1$weight and Diet3$weight  
t = -3.4293, df = 16.408, p-value = 0.003337  
alternative hypothesis: true difference in means is not equal to 0  
95 percent confidence interval:  
-149.64644 -35.45356  
sample estimates:  
mean of x mean of y  
177.75 270.30
```

##### Working with paired data is another common statistical activity.

```
# The "ChickWeight" data will be used to illustrate how the weight gain  
# from day 20 to 21 may be analyzed. Use the following code to prepare  
# pre- and post-data from Diet == "3" for analysis.
```

```
#``{r test3paired}
```

```
# load "ChickWeight" dataset
```

```
data(ChickWeight)
```

```
# Create T | F vector indicating observations with Diet == "3"  
index <- ChickWeight$Diet == "3"  
  
# Create vector of "weight" for observations where Diet == "3" and Time == 20  
pre <- subset(ChickWeight[index, ], Time == 20, select = weight)$weight
```

```

# Create vector of "weight" for observations where Diet == "3" and Time == 21
post <- subset(ChickWeight[index, ], Time == 21, select = weight)$weight

# The pre and post values are paired, each pair corresponding to an individual chick.
cbind(pre, post)

#(3)(c) (3 points) Present a scatterplot of the variable "post" as a function of the
# variable "pre". Include a diagonal line with zero intercept and slope equal to one.

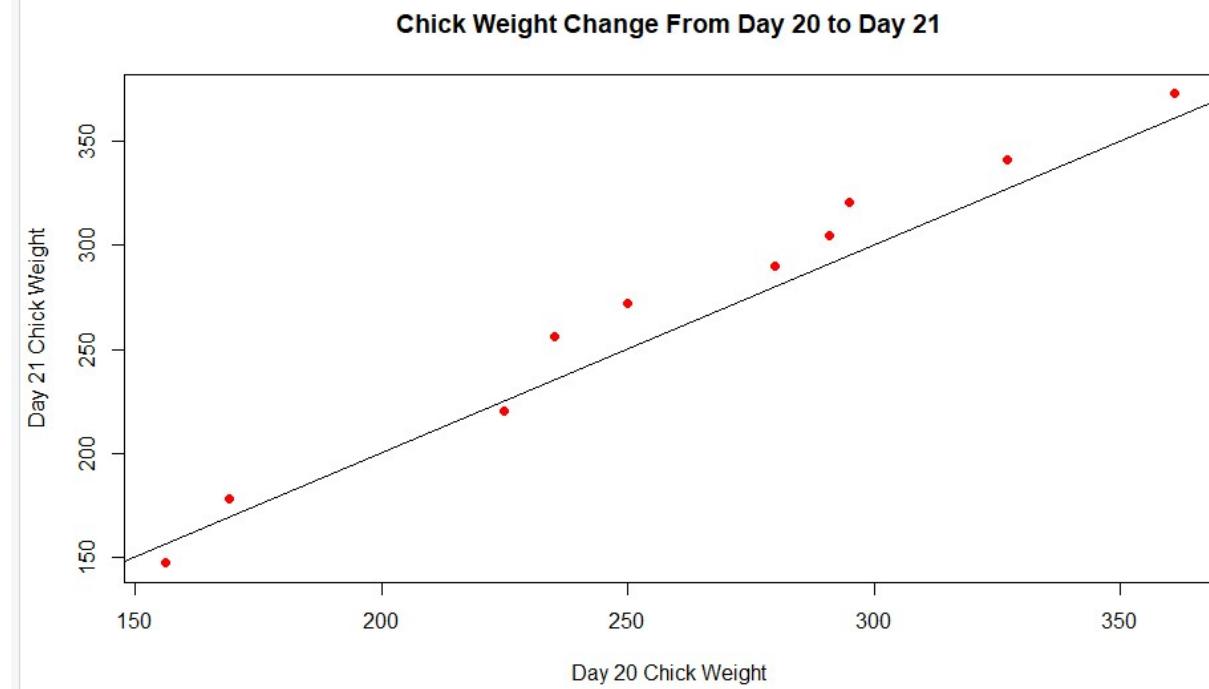
# Title and label the variables in this scatterplot.

plot(pre, post, xlab="Day 20 Chick Weight",
      ylab="Day 21 Chick Weight",
      main = "Chick Weight Change From Day 20 to Day 21",
      pch=16, col="red")

abline(a=0,b=1)

output
#####
> plot(pre, post, xlab="Day 20 chick weight",
+       ylab="Day 21 chick weight",
+       main = "Chick weight Change From Day 20 to Day 21",
+       pch=16, col="red")
> abline(a=0,b=1)

```



```

#(3)(d) (6 points) Calculate and present a one-sided, 95% confidence interval for the
# average weight gain from day 20 to day 21. Write the code for the paired t-test and
# for determination of the confidence interval endpoints. **Do not use *t.test()**,
# although you may check your answers using this function. Present the resulting
# test statistic value, critical value, p-value and confidence interval.

#``{r test3d}

mu<-mean(pre - post)

n<-length(pre)

crit_t<-qt(p=.950,df=n-1)

se<-sd(pre-post)/sqrt(n)

t<-mu/se

list(t.statistic = t, critical.t = crit_t,
      p.value = 2 * pt(q = t, df = n - 1, lower.tail = FALSE),
      conf.interval = c(mu - crit_t * se, mu + crit_t * se))

t.test(x = pre, y = post, paired = TRUE)

output
#####
> mu<-mean(pre - post)
> n<-length(pre)
> crit_t<-qt(p=.950,df=n-1)
> se<-sd(pre-post)/sqrt(n)
> t<-mu/se
>
> list(t.statistic = t, critical.t = crit_t,
+       p.value = 2 * pt(q = t, df = n - 1, lower.tail = FALSE),
+       conf.interval = c(mu - crit_t * se, mu + crit_t * se))
$t.statistic
[1] -3.225267

$critical.t
[1] 1.833113

$p.value
[1] 1.989599

$conf.interval
[1] -17.879304 -4.920696

>

```

```

> t.test(x = pre, y = post, paired = TRUE)

Paired t-test

data: pre and post
t = -3.2253, df = 9, p-value = 0.0104
alternative hypothesis: true difference in means is not equal to 0
95 percent confidence interval:
-19.3958 -3.4042
sample estimates:
mean of the differences
-11.4

```

#(4)(a) (3 points) Using Nile River flow data and the "moments" package, calculate  
# skewness and kurtosis. Present a QQ plot and boxplot of the flow data side-by-side  
# using \*qqnorm()\*, \*qqline()\* and \*boxplot()\*; \*par(mfrow = c(1, 2))\* may be used  
# to locate the plots side-by-side. Add features to these displays as you choose.

```

#``{r test4a}

library(moments)

str(Nile)

summary(Nile)

skewness(Nile)

kurtosis(Nile)

par(mfrow=c(1,2))

boxplot(Nile, col = 'red', main = 'Nile River Flows 1871 to 1970' )

qqnorm(Nile, col = 'red', pch=16, main = 'Nile River Flows 1871 to 1970' )

qqline(Nile)

par(mfrow=c(1,1))

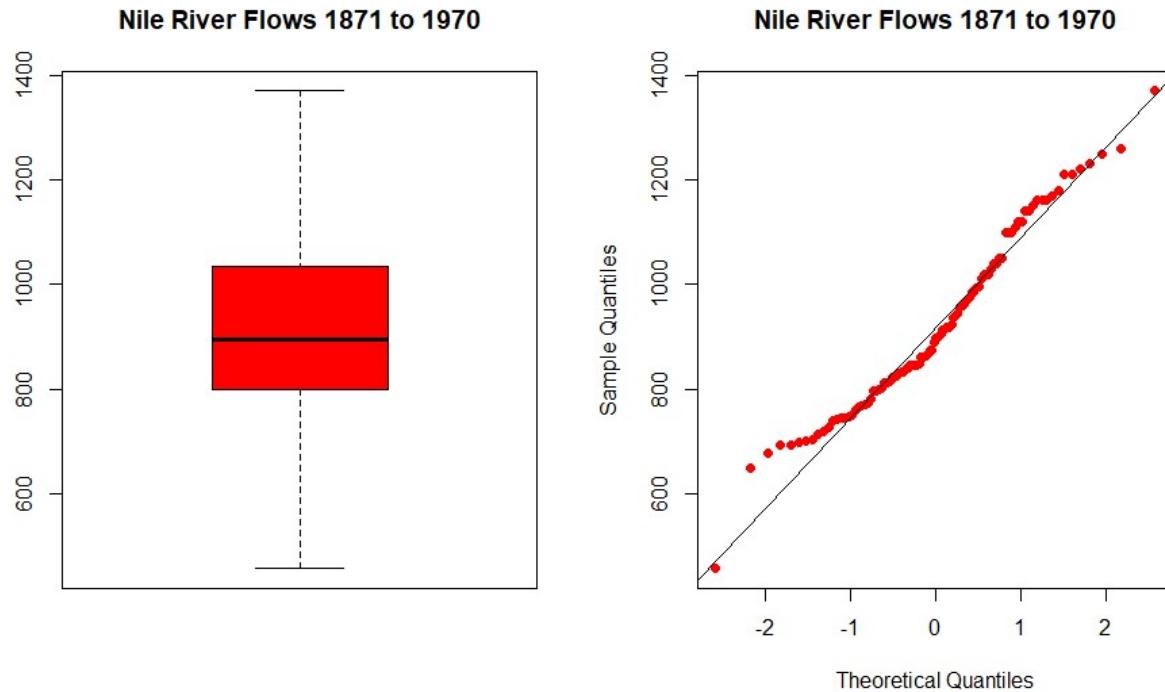
output
#####
> library(moments)
> str(Nile)
Time-Series [1:100] from 1871 to 1970: 1120 1160 963 1210 1160 ...
1370 1140 ...
> summary(Nile)
   Min. 1st Qu. Median     Mean 3rd Qu.    Max.
456.0  798.5  893.5  919.4 1032.5 1370.0

```

```

> skewness(Nile)
[1] 0.3223697
> kurtosis(Nile)
[1] 2.695093
> par(mfrow=c(1,2))
> boxplot(Nile, col = 'red', main = 'Nile River Flows 1871 to 1970' )
> qqnorm(Nile, col = 'red', pch=16, main = 'Nile River Flows 1871 to 1970' )
> qqline(Nile)
> par(mfrow=c(1,1))

```



```

#(4)(b) (6 points) Using *set.seed(124)* and the Nile data, generate 1000 random
# samples of size n = 16, with replacement. For each sample drawn, calculate and
# store the sample mean. This can be done with a for-loop and use of the *sample()*
# function. Label the resulting 1000 mean values as "sample1". **Repeat these steps
# using *set.seed(127)* - a different "seed" - and samples of size n = 64.** Label
# these 1000 mean values as "sample2". Compute and present the means, sample standard
# deviations and sample variances for "sample1" and "sample2" in a table with the first
# row for "sample1", the second row for "sample2" and the columns labeled for each statistic.

#``{r test4b}

```

```

set.seed(124)
sample1<-numeric(1000)
for (i in 1:1000) {
  a<-sample(y, 16, replace = T)
  sample1[i]<- mean(a)
}
set.seed(127)
sample2<-matrix(, nrow = 1000, ncol=3)
sample2<-numeric(1000)
for (i in 1:1000) {
  a<-sample(y, 64, replace = T)
  sample2[i]<- mean(a)
}
(samp <- cbind(mean(sample1),sd(sample1),var(sample1)))
(samp <- rbind(samp,cbind(mean(sample2),sd(sample2), var(sample2))))
rownames(samp) <- c("sample1","sample2")
colnames(samp) <- c("Mean", "Std", "Variance")
samp

```

#### OUTPUT

---

```

> set.seed(124)
> sample1<-numeric(1000)
> for (i in 1:1000) {
+   a<-sample(y, 16, replace = T)
+   sample1[i]<- mean(a)
+ }
>
> set.seed(127)
> sample2<-matrix(, nrow = 1000, ncol=3)
> sample2<-numeric(1000)
> for (i in 1:1000) {
+   a<-sample(y, 64, replace = T)
+   sample2[i]<- mean(a)
+ }
>
> (samp <- cbind(mean(sample1),sd(sample1),var(sample1)))
     [,1]     [,2]     [,3]
[1,] 0.1878658 0.0469961 0.002208633
> (samp <- rbind(samp,cbind(mean(sample2),sd(sample2), var(sample2))))
     [,1]     [,2]     [,3]
[1,] 0.1878658 0.04699610 0.002208633

```

```

[2,] 0.1889880 0.02346896 0.000550792
> rownames(samp) <- c("sample1", "sample2")
> colnames(samp) <- c("Mean", "Std", "Variance")
>
> samp
      Mean       Std   Variance
sample1 0.1878658 0.04699610 0.002208633
sample2 0.1889880 0.02346896 0.000550792

#(4)(c) (6 points) Present side-by-side histograms of "sample1" and "sample2" with the
# normal density curve superimposed. To prepare comparable histograms, it will be
# necessary to use "freq = FALSE" and to maintain the same x-axis with
# "xlim = c(750, 1050)", and the same y-axis with "ylim = c(0, 0.025)."
# **To superimpose separate density functions, you will need to use the mean and
# standard deviation for each "sample" - each histogram - separately.**
#``{r test4c}

par(mfrow=c(1,2))

x<-seq(from = 700, to = 1100, by = 1)

hist(sample1, freq = FALSE, col = "darkred", xlab = "Flow",
      main = "Histogram of Nile River\nSample 1 Flows 1871 to 1970"
      , ylim = c(0, 15),
      xlim = c(0, 0.25), plot = TRUE)

curve(dnorm(x, mean = samp[1,1], sd = samp[1,2]), col = "green", lwd = 2, add = TRUE)

hist(sample2, freq = FALSE, col = "darkred", xlab = "Flow",
      main = "Histogram of Nile River\nSample 2 Flows 1871 to 1970"
      , ylim = c(0, 15),
      xlim = c(0, 0.25), plot = TRUE)

curve(dnorm(x, mean = samp[2,1], sd = samp[2,2]), col = "green", lwd = 2, add = TRUE)

par(mfrow=c(1,1))

output

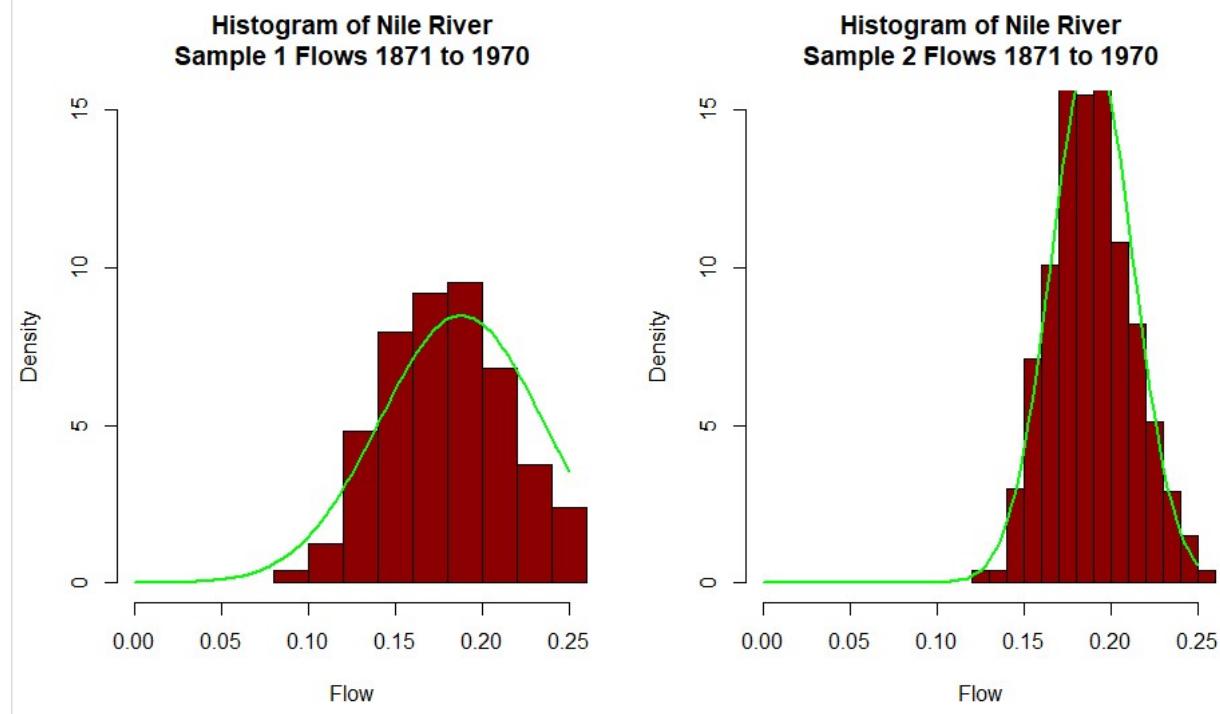
> par(mfrow=c(1,2))
> x<-seq(from = 700, to = 1100, by = 1)
> hist(sample1, freq = FALSE, col = "darkred", xlab = "Flow",
+       main = "Histogram of Nile River\nSample 1 Flows 1871 to 1970"
+       , ylim = c(0, 15),
+       xlim = c(0, 0.25), plot = TRUE)

```

```

> curve(dnorm(x, mean = samp[1,1], sd = samp[1,2]), col = "green", lwd = 2, add = TRUE)
> hist(sample2, freq = FALSE, col = "darkred", xlab = "Flow",
+      main = "Histogram of Nile River\nSample 2 Flows 1871 to 1970"
+ , ylim = c(0, 15),
+ xlim = c(0, 0.25), plot = TRUE)
> curve(dnorm(x, mean = samp[2,1], sd = samp[2,2]), col = "green", lwd = 2, add = TRUE)
> par(mfrow=c(1,1))

```



## ##### Section 5: (15 points) #####

##### (5) This problem deals with contingency table analysis. This is an example  
# of categorical data analysis (see Kabacoff, pp. 145-151). The "warpbreaks" dataset  
# gives the number of warp breaks per loom, where a loom corresponds to a fixed length  
# of yarn. There are 54 observations on 3 variables: breaks (numeric, the number of breaks),  
# wool (factor, type of wool: A or B), and tension (factor, low L, medium M and high H).  
# These data have been studied and used for example elsewhere. For the purposes of this problem,  
# we will focus on the relationship between breaks and tension using contingency table analysis.

```
# (5)(a)(4.5 points) warpbreaks is part of the "datasets" package and may be loaded via  
# *data(warpbreaks)*. Load "warpbreaks" and present the structure using *str()*.  
  
# Calculate the median number of breaks for the entire dataset, disregarding "tension" and  
# "wool". Define this median value as "median_breaks". Present a histogram of the number of  
# breaks with the location of the median indicated.
```

```
# Create a new variable "number" as follows: for each value of "breaks", classify the number  
# of breaks as either strictly below "median_breaks", or the alternative. Convert the  
# "above" | "below" classifications to a factor, and combine with the dataset warpbreaks.  
  
# Present a summary of the augmented dataset using *summary()**. Present a contingency  
# table of the frequency of breaks using the two variables "tension" and "number".  
  
# There should be six cells in this table.
```

```
#``{r test5a}
```

```
data(warpbreaks)  
warp<-warpbreaks  
str(med)
```

```
(median_breaks <- median(warp$breaks))
```

```
hist(warp$breaks, main = paste("Histogram of Breaks with Median Value\n Blue vertical line at ",  
                               median_breaks, "Breaks"))  
abline(v= median_breaks, col="blue")
```

```
warp$P50.loc <- factor(warp)  
warp$P50.loc <- factor(warp$breaks>median_breaks)  
levels(warp$P50.loc <- c("below", "above"))
```

```

summary(warp)

(warptab <- table(warp$tension,warp$P50.loc))

#```

data(warpbreaks)

warp<-warpbreaks

str(med)

(median_breaks <- median(warp$breaks))

hist(warp$breaks, main = paste("Histogram of Breaks with Median Value\n Blue vertical line at ", median_breaks, "Breaks"))

abline(v= median_breaks, col="blue")

warp$P50.loc <- factor(warp)

warp$P50.loc <- factor(warp$breaks>median_breaks)

levels(warp$P50.loc <- c("below", "above"))

summary(warp)

(warptab <- table(warp$tension,warp$P50.loc))

```

Output

---

```

> data(warpbreaks)
> warp<-warpbreaks
> str(med)
'data.frame': 54 obs. of 4 variables:
 $ breaks : num 26 30 54 25 70 52 51 26 67 18 ...
 $ wool    : Factor w/ 2 levels "A","B": 1 1 1 1 1 1 1 1 1 1 ...
 $ tension: Factor w/ 3 levels "L","M","H": 1 1 1 1 1 1 1 1 1 2 ...
 $ P50.loc: chr "below" "above" "below" "above" ...
>
> (median_breaks <- median(warp$breaks))
[1] 26
>
> hist(warp$breaks, main = paste("Histogram of Breaks with Median Value\n Blue vertical line at ", median_breaks, "Breaks"))
+ abline(v= median_breaks, col="blue")
>
> warp$P50.loc <- factor(warp)
> warp$P50.loc <- factor(warp$breaks>median_breaks)
> levels(warp$P50.loc <- c("below", "above"))
NULL

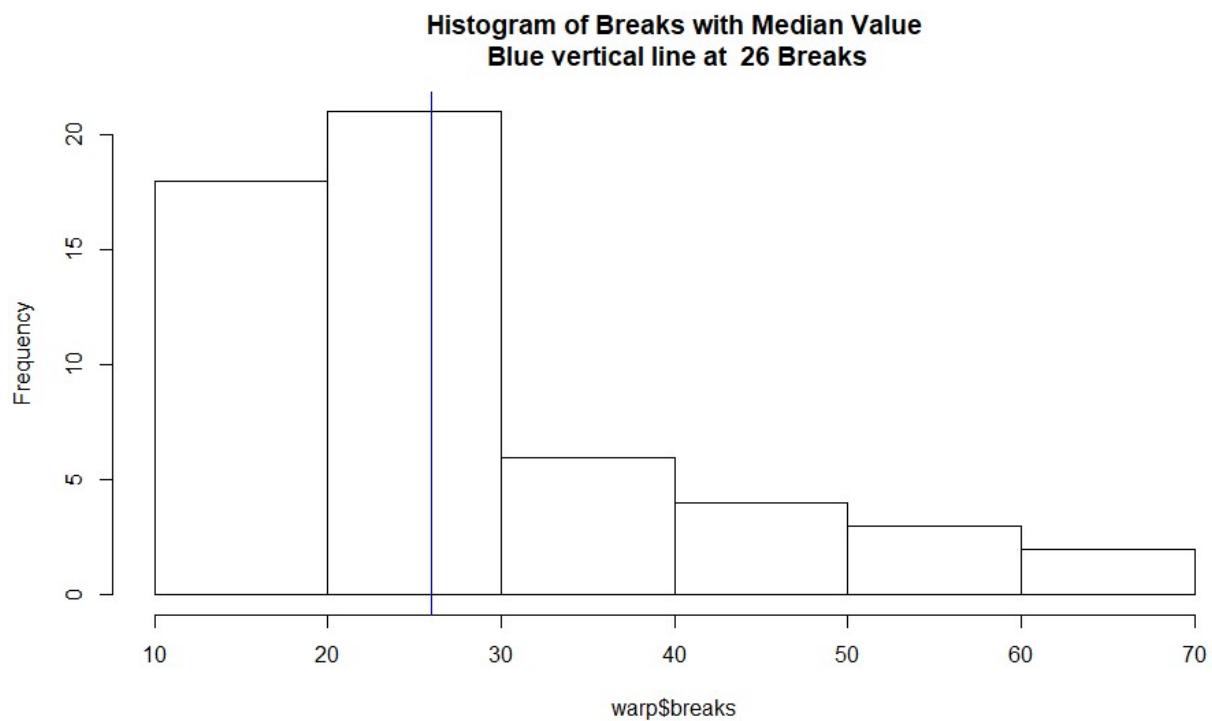
```

```

> summary(warp)
  breaks      wool      tension    P50.loc
Min.   :10.00   A:27      L:18       Length:54
1st Qu.:18.25  B:27      M:18       Class  :character
Median  :26.00          H:18       Mode   :character
Mean    :28.15
3rd Qu.:34.00
Max.   :70.00
>
> (warptab <- table(warp$tension,warp$P50.loc))

  above below
L     9     9
M     9     9
H     9     9

```



```

#(5)(b)(3 points) Using the table constructed in (5)(a), test at the 5% level
# the null hypothesis of independence using the uncorrected *chisq.test()*
# (Black, Business Statistics, Section 16.2). Show the results of this test and state your conclusions.

#``{r test5b}

chisq.test(warptab, correct = FALSE)

output
#####
> chisq.test(warptab, correct = FALSE)

Pearson's Chi-squared test

data: warptab
X-squared = 0, df = 2, p-value = 1

```

The observed value of chi squared 0 is less than the critical value 5.9915 so the null hypothesis tension is independent of number is accepted.

```

#(5)(c) (7.5 points) Write a function that computes the uncorrected
# Pearson Chi-squared statistic. Apply your function to the table from (5)(a).

# You should be able to duplicate the X-squared value (chi-squared) and *p*-value. Present both.

```

```

# Shown below are examples of the type of function required. These examples will have
# to be modified to accomodate the table generated in (5)(a).

```

```

#``{r test5c}

chi <- function(x) {

  # To be used with 2x2 contingency tables that have margins added.

  # Expected values are calculated.

  e11 <- x[4,1]*x[1,3]/x[4,3]
  e12 <- x[4,2]*x[1,3]/x[4,3]

```

```

e21 <- x[4,1]*x[2,3]/x[4,3]
e22 <- x[4,2]*x[2,3]/x[4,3]
e31 <- x[4,1]*x[3,3]/x[4,3]
e32 <- x[4,2]*x[3,3]/x[4,3]

# Value of chi square statistic is calculated.

chisqStat <- (x[1,1] - e11)^2/e11 + (x[1,2] - e12)^2/e12 +
(x[2,1] - e21)^2/e21 + (x[2,2] - e22)^2/e22 +
(x[3,1] - e31)^2/e31 + (x[3,2] - e32)^2/e32

return(list("chi-squared" = chisqStat,
           "p-value" = pchisq(chisqStat, 1, lower.tail = F)))
}

```

```

chisqfun <- function(t) {

  x <- addmargins(t)

  e <- matrix(0, nrow = nrow(t), ncol = ncol(t), byrow = T)
  r <- matrix(0, nrow = nrow(t), ncol = ncol(t), byrow = T)

  for (i in 1:2) {
    for (j in 1:2) {
      e[i,j] = x[nrow(x),j] * x[i,ncol(x)]/x[nrow(x), ncol(x)]
      r[i,j] = ((x[i,j] - e[i,j])^2)/e[i,j]
    }
  }

  chi <- sum(r)

  xdf <- nrow(t) - 2

  pv <- pchisq(chi, df = xdf, lower.tail = FALSE)

  return(cat("Pearson's Chi-squared test \n","Chi sq: ", chi, ";",
            "Degree of Freedom : ", xdf, " ; P-value : ", pv))
}

```

```
(Y <- matrix(warptab, nrow = 3, byrow = F))
(A <- cbind(Y, apply(Y, 1, sum)))
(B <- rbind(A, apply(A, 2, sum)))
```

chi(B)

chisqfun(B)

output

---

```
> chi <- function(x) {
+   # To be used with 2x2 contingency tables that have margins added.
+   # Expected values are calculated.
+   e11 <- x[4,1]*x[1,3]/x[4,3]
+   e12 <- x[4,2]*x[1,3]/x[4,3]
+   e21 <- x[4,1]*x[2,3]/x[4,3]
+   e22 <- x[4,2]*x[2,3]/x[4,3]
+   e31 <- x[4,1]*x[3,3]/x[4,3]
+   e32 <- x[4,2]*x[3,3]/x[4,3]
+   # Value of chi square statistic is calculated.
+   chisqstat <- (x[1,1] - e11)^2/e11 + (x[1,2] - e12)^2/e12 +
+     (x[2,1] - e21)^2/e21 + (x[2,2] - e22)^2/e22 +
+     (x[3,1] - e31)^2/e31 + (x[3,2] - e32)^2/e32
+   return(list("chi-squared" = chisqstat,
+             "p-value" = pchisq(chisqstat, 1, lower.tail = F)))
+ }
>
> chisqfun <- function(t) {
+   x <- addmargins(t)
+   e <- matrix(0, nrow = nrow(t), ncol = ncol(t), byrow = T)
+   r <- matrix(0, nrow = nrow(t), ncol = ncol(t), byrow = T)
+   for (i in 1:2) {
+     for (j in 1:2) {
+       e[i,j] = x[nrow(x),j] * x[i,ncol(x)]/x[nrow(x), ncol(x)]
+       r[i,j] = ((x[i,j] - e[i,j])^2)/e[i,j]
+     }
+   }
+   chi <- sum(r)
+   xdf <- nrow(t) - 2
+   pv <- pchisq(chi, df = xdf, lower.tail = FALSE)
+   return(cat("Pearson's Chi-squared test \n", "Chi sq: ", chi, ";",
+             "Degree of Freedom :", xdf, " ; P-value : ", pv))
+ }
>
> (Y <- matrix(warptab, nrow = 3, byrow = F))
 [,1] [,2]
[1,]    9    9
[2,]    9    9
[3,]    9    9
> (A <- cbind(Y, apply(Y, 1, sum)))
 [,1] [,2] [,3]
```

```
[1,]    9    9   18
[2,]    9    9   18
[3,]    9    9   18
> (B <- rbind(A, apply(A, 2, sum)))
 [,1] [,2] [,3]
[1,]    9    9   18
[2,]    9    9   18
[3,]    9    9   18
[4,]   27   27   54
>
> chi(B)
$`chi-squared`
[1] 0

$`p-value`
[1] 1

> chisqfun(B)
Pearson's Chi-squared test
Chi sq: 0 ;
Degree of Freedom : 2 ; P-value : 1
```